

**For both optimizers, BF16 achieves far higher accuracy per unit of KL,
but requires 10x more training steps.**

GSM8K test set · bfloat16 inference · greedy decoding

