

L_0 sparsity and k_{90} are both sensitive to precision —
under mixed precision, where these metrics are uncontaminated by truncation effects,
AdamW's updates are $\sim 3\times$ more spectrally concentrated than Muon's

