


BF16 training appears to diverge far less in the forward KL space while achieving comparable accuracy on GSM8K format learning (Qwen2-1.5B-Instruct)

(a) Validation Accuracy vs Forward KL



(b) Validation Accuracy vs Training Compute

