

Calibration: SFT vs GRPO vs RS (~1.1M tokens, FP32 logprobs, greedy decoding)

Optimizer	Objective	Accuracy	Parsable	ECE	Mean Conf.
AdamW	SFT	50.0%	98.9%	0.4881	0.9939
Muon	SFT	51.0%	99.2%	0.4786	0.9931
AdamW	GRPO	61.4%	98.4%	0.3718	0.9877
Muon	GRPO	59.3%	95.8%	0.3776	0.9963
AdamW	RS	59.1%	98.8%	0.3983	0.9961
Muon	RS	61.9%	98.3%	0.3675	0.9979